

IST 659: Server inventory and network management system

Final Project Report

Instructor: Prof. Chad Harper

Team:
Sharat Sripada

TABLE OF CONTENTS

INTRODUCTION	3
PROBLEM STATEMENT	3
BUSINESS RULES AND ENTITY RELATIONSHIP DIAGRAM (ERD)	3
ABSTRACT	3
MODELING ENTITIES AND ATTRIBUTES	4
ENTITY RELATIONSHIP DIAGRAM (ERD)	5
LOGICAL MODELING AND ENHANCED ENTITY RELATIONSHIP DIAGRAM (E-ERD)	5
NORMALIZATION	6
IMPLEMENTATION – DIVE INTO MY FIRST FLASK APPLICATION	7
EXPLORING DATABASE CREATE, READ, UPDATE, DELETE OPERATIONS	8
CREATE RDBMS DATABASE AND TABLES	8
API GATEWAY – @APP.ROUTE DECORATORS	9
POPULATING DATA – A FULLY FUNCTIONING DATABASE	10
PYTHON UTILITY	10
FORKING FUNCTION – BRIDGE BETWEEN API AND DATABASE	11
SQL – UNDER THE HOOD DATA PERSISTENCE	11
ANALYTICS	11
SQL – BEYOND NATIVE CRUD	12
DASHBOARD	12
CONCLUSION AND FUTURE WORK	13
FUTURE WORK	13
APPENDIX	14
FRONT-END CODE FOR DASHBOARD (CHART.HTML)	14

Introduction

Server inventory and network management at scale has been a pain point for our Engineering Team. With hundreds of servers, switches and switch ports to manage across several production environments it sets a nearly perfect use-case to brace an RDBMS at the backend to store large records of structured data while presenting a simplistic frontend.

Through the knowledge acquired in IST-659 and guidance from Prof. Chad Harper the intent is to design a solution that will help greatly ease the day in a life of a developer.

Problem Statement

A production (aka prod.) environment typically comprises several servers and upstream switches (CLOS fabric comprising first-hop switches termed LEAF and a SPINE aggregate). Developers are building these prod. environments typically doing the following:

- Add servers to prod. to increase capacity
(OR)
Remove servers from prod. when servicing hardware faults or deprecating servers (due to end of life/expiry of service contracts)
- Configuring upstream switches to facilitate workloads or virtual machines (VM) on prod. servers to communicate with each other across the datacenter
- Setting up NFS mounts etc.
- Shuffling servers between prod. environments

Clearly, this is a cumbersome process proving detrimental to Engineering productivity. The endeavor is to automate the above steps end-to-end offering simplistic workflows via a Graphical User interface (GUI) or Application Programmable Interface (API) and utilizing an RDBMS at the backend.

Business rules and Entity Relationship Diagram (ERD)

Abstract

Consider the following narrative explaining the Business rules and scope:

Engineers run their code or features at scale building up prod. environments which would comprise servers and associative switches to facilitate communication.

Each prod. environment will have a dedicated VMWare vCenter managing workloads (Outer: to manage physical servers and Inner: to manage nested hypervisors), Runner virtual-machine (to stage, compile and execute code), an NFS mount-point which includes an ip-address and path, more than ONE dedicated virtual-LAN/VLAN IDs isolating Management, Overlay/data-path and NFS traffic.

Prod. environments are backed by ONE or more servers running VMWare ESX hypervisor and would comprise data related to server-manufacturer, a single ip-address assigned to the server to access it over the Management network, vmkernel ip-address to access NFS, memory and CPU capacities of the server and a parameter indicating whether the server was in-use. If a server were not to be in use it could either be deprecated or temporarily out-of-service servicing a hardware fault/upgrade.

Finally, all servers are managed by TWO upstream Layer-2 switches. The switch would comprise data regarding the manufacturer/vendor (this may also be needed to adjust the switch CLI syntax to automate networking actions), port, port-type, switch ip-address, switch type (Leaf/Spine) and finally more than one virtual-LAN/VLAN IDs isolating Management, Overlay/data-path and NFS traffic mapped to each port.

NOTE

Deprecated servers shall be cleaned up from the prod. environment correspondingly also cleaning up any network related references to switch or switch port.

Modeling Entities and Attributes

Based on the abstract following are the entities, attributes & datatypes:

Entities	Attributes	Attribute property	Datatype
Production	Prod. Name	Primary Key	varchar
	NFSIPAddress	Required	varchar
	NFSShare (path)	Required	varchar
	OuterVCIPAddress	Required, Unique	varchar
	InnerVCIPAddress	Required, Unique	varchar
	{VLANs}	Required, Multiple	Int
Server	Server Name	Primary Key	varchar
	Manufacturer	Required	varchar
	IPAddress	Required, Unique	varchar
	VMKernelIPAddress	Required, Unique	varchar
	MemoryInGB	Required	float
	CPUInCores	Required	int
	InUse	Required	varchar
	{VLANs}	Required, Multiple	int
Switch	Switch Name	Primary Key	varchar
	Manufacturer	Required	varchar
	Port	Required	varchar
	PortType	Required	varchar
	IPAddress	Required, Unique	varchar
	SwitchType	Required	varchar
FaultTicket	FaultID	Primary Key	varchar
	Description	Optional	varchar
	Severity	Optional	varchar

Entity Relationship Diagram (ERD)

The entities, attributes and relationships can be represented using an ERD. The relationships between entities is detailed as:

- **Prod. and Server entities**

A prod. environment will comprise 1 or many servers

(AND)

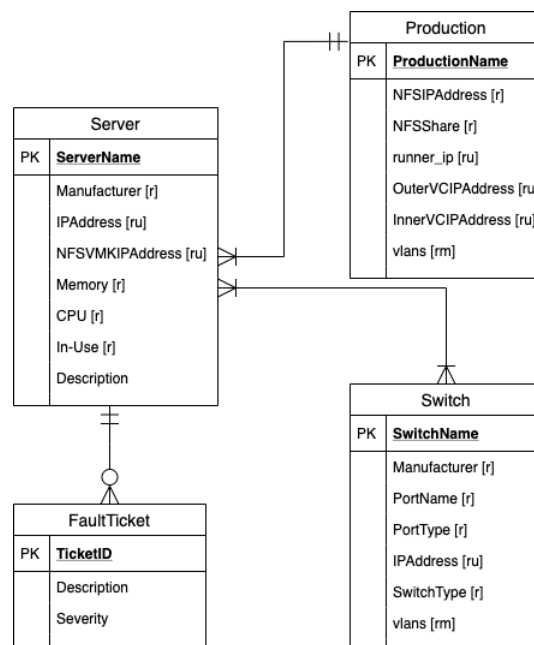
A server can belong to ONE and ONLY ONE prod. environment, hence a ONE-MANY relationship

- **Server and Switch entities**

A server would have its physical ports on Network Interface Card (NIC) connected to TWO (or multiple) switches.

(AND)

Likewise, a switch can comprise several servers connected to it, hence a MANY-MANY relationship.



Sharat Sripada - Project
(ERD)

Figure-1: Entity Relationship Diagram

Logical Modeling and Enhanced Entity Relationship Diagram (E-ERD)

Translating the ERD in Figure-1 to a logical diagram, we do the following:

- Map entities to tables
- In some cases, prod. name or switch name is too long and sometimes cannot be guaranteed uniqueness so, use surrogate keys instead (int identity datatype)
- All attributes are correspondingly mapped to column names following the best-practices for naming conventions (all lower case & separate with underscore for spaces)
- Map relationships between entities as relationships between tables

- Map the MANY-MANY relationship between Server and Switch entities to a MANY-ONE/ONE-MANY relationship using a bridge or associative table called server_switch_info
- Vlan's are multi-valued and comprise other non-key attributes like Vlan type indicating Management, Overlay or NFS networks. This is correspondingly mapped to the table belongs and vlan_table.

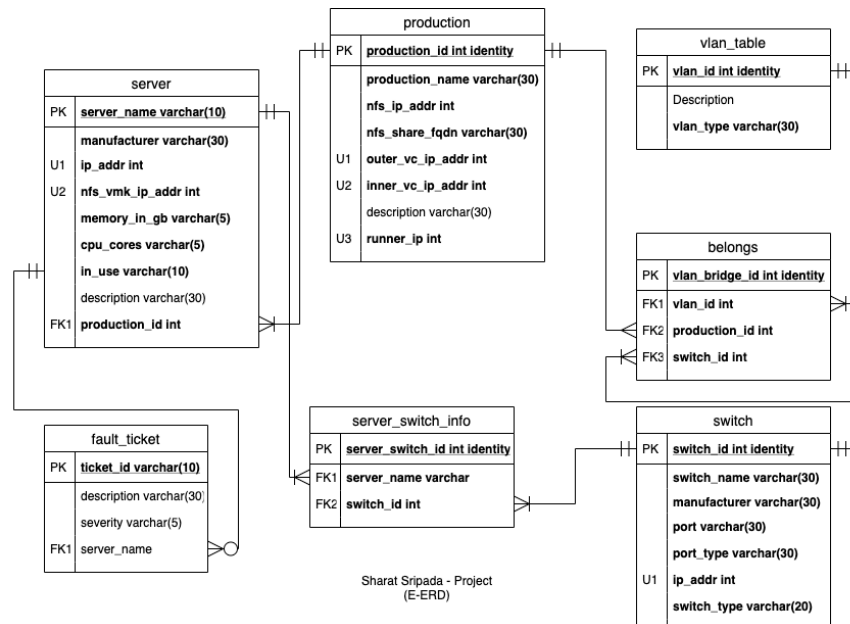


Figure-2: Enhanced Entity Relationship Diagram

Normalization

After placing real data from THREE Prod. environments and running the tables through the normalization steps (1NF -> 2NF -> 3NF process), saw the need for a few changes.

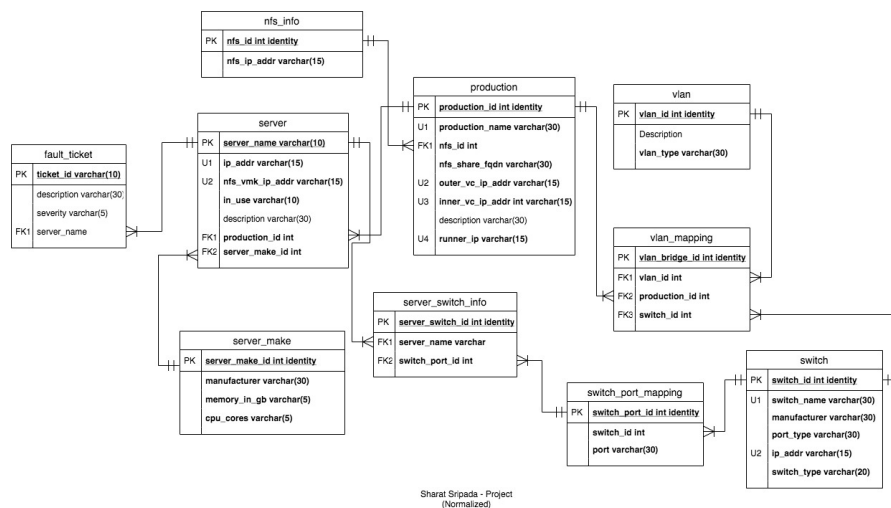


Figure-3: 3NF tables after normalization

Summary of changes:

- Total number of tables increased from SEVEN to TEN. Newly added tables are `nfs_info`, `switch_port_mapping` and `server_make`
- Changed type for ip-address column-names from `int` to `varchar(15)`
- Change name of table 'belongs' to 'vlan_mapping'

Implementation – Dive into my first Flask application

For the implementation, Flask - a web application framework in Python was chosen. The framework along with certain libraries offers a rich platform to develop an end-to-end solution comprising:

- Graphical User Interface (GUI) – Using the `@app.route()` decorator build HTML with embedded JavaScript/AJAX for dynamic front-end objects/charts
- API framework – Using the `@app.route()` decorator to build POST, GET, PATCH, DELETE API methods. This would then correspondingly translate to CREATE, SELECT, UPDATE and DELETE SQL operations
- RDBMS/SQL database - Sqlite3 libraries for create and administer databases

The code is organized and maintained in the following directory structure:




Figure-4: Organizing files (Legend: Orange fills indicate directories)

The application itself is developed within `app.py` under the parent directory `/inventory` which when running would listen on port `*:5000` per this in the code:

```
/inventory[687](my_flask_app)# less app.py.
```

```
558 app.run(host='0.0.0.0')
```

Other directories and files:

- `/db` comprise the sqlite database files (`.db` files)
- `/raw_data` comprise data in csv format translated from excel files (representing database tables in 3NF).
 - `populate_db.py` will read data and populate corresponding tables
- `/templates` will comprise all essential front-end related html/java-script files

Here's a detailed dump of directories and files under /inventory:

```
# ls *
app.py  populate_db.py

db:
test.db  testinfo.db  testinfra.db

raw_data:
fault_ticket  nfs_info  prod  server  server_make  server_switch_info  switch
switch_port_mapping  vlan  vlan_mapping

static:

templates:
chart.html  index.html
```

NOTE:

The files under raw_data correspond 1:1 to tables in the database.

Exploring database Create, Read, Update, Delete operations

With all necessary constructs in place this section now encompasses details of everything that went into building the database, corresponding tables and performing the Create, Read, Update, Delete operations, commonly hence referred as CRUD ops.

NOTE: At the time of putting this report in place, only Create and Read were implemented. See section Future Work for further details regarding pending tasks.

Create RDBMS database and tables

Using the Python sqlite3 library, capturing an excerpt of how the tables were created using SQL in *db/testinfra.db* database. Since table creation is within the `__init__` constructor, the corresponding CREATE SQL commands kick in on `class SqlLibTestInfra()` instantiation. Checks are in place to walk all tables not comprising 'sqlite' (see LINES 19-20) prior, to prevent redundant work.

Table creation is in a specific sequence adhering to CONSTRAINTs or references to FOREIGN_KEY attributes.

Below is the corresponding code:


```

>> 13 class SqlLibTestInfra(object):
>> 14     def __init__(self):
15         self.conn = sqlite3.connect('db/testinfra.db', check_same_thread=False)
16         _user_tables = ('nfs_info', 'production', 'server', 'server_make',
17                         'fault_ticket', 'switch_port_mapping', 'switch',
18                         'server_switch_info', 'vlan_mapping', 'vlan')
19         _res = self.conn.execute('select name from sqlite_master where \
20                                 type=\'table\' and name NOT LIKE \'sqlite_%\';')
21         db_tables = _res.fetchall()
22         for user_table in _user_tables:
23             table_present = False
24             for db_table in db_tables:
25                 if user_table == db_table[0]:
26                     table_present = True
27                     break
28             if not table_present:
29                 if user_table == 'nfs_info':
30                     self.conn.execute('create table nfs_info ( \
31                                     nfs_id int PRIMARY KEY, \
32                                     nfs_ip_addr varchar(15) not null);')
33                 if user_table == 'production':
34                     self.conn.execute('create table production ( \
35                                     production_id int PRIMARY KEY, \
36                                     production_name varchar(2) not null, \
37                                     nfs_share_fqdn varchar(30) not null, \
38                                     outer_vc_ip varchar(15) not null, \
39                                     inner_vc_ip varchar(15) not null, \
40                                     description varchar(15) not null, \
41                                     runner_ip varchar(15) not null, \
42                                     nfs_id int not null, \
43                                     FOREIGN KEY (nfs_id) \
44                                     REFERENCES nfs_info(nfs_id));')
45                 if user_table == 'server':
46                     self.conn.execute('create table server ( \
47                                     server_name varchar(30) PRIMARY KEY, \
48                                     server_make_id int, \
49                                     ip_addr varchar(15), \
50                                     nfs_vmk_ip_addr varchar(15), \
51                                     in_use varchar(10), \
52                                     description varchar(20), \
53                                     production_id int not null, \
54                                     FOREIGN KEY (production_id) \
55                                     REFERENCES production(production_id));')

```

Figure-5: Table creation code using SQLite libraries/commands

API gateway – @app.route decorators

Flask has `@app.route` decorators to build very comprehensive API gateways. Below is a proponent of its usage to cater to common POST and GET methods – `common_db_crud_ops` (more of this in the next section) then forks the requests to SQL INSERT or SELECT statements respectively.

APIs would eventually be extended to REST methods PATCH and DELETE that correspond to SQL UPDATE and DELETE statements respectively.

```

476 # Server/switch inventory - apis
477 @app.route('/api/v1/production/', methods=['POST', 'GET'])
>> 478 def production():
~ 479     return common_db_crud_ops('production', request)
480
481 @app.route('/api/v1/nfs_info/', methods=['POST', 'GET'])
>> 482 def nfs_info():
~ 483     return common_db_crud_ops('nfs_info', request)
484
485 @app.route('/api/v1/server/', methods=['POST', 'GET'])
>> 486 def server():
~ 487     return common_db_crud_ops('server', request)
488
489 @app.route('/api/v1/server_make/', methods=['POST', 'GET'])
>> 490 def server_make():
~ 491     return common_db_crud_ops('server_make', request)
492
493 @app.route('/api/v1/fault_ticket/', methods=['POST', 'GET'])
>> 494 def fault_ticket():
~ 495     return common_db_crud_ops('fault_ticket', request)

```

Figure-6: APIs implemented to insert, get data

Populating data – A fully functioning database

With API gateway and database in place, we start to populate the tables with real data.

Explored mockaroo.com but data corresponding to some tables or columns was extremely tailored to our use-case. So, all data is real - gathered from a few prod. environments (this is a tedious/manual process and is still ongoing). It is then organized in fully compliant 3NF equivalent excel sheets before converting and exporting them in .csv format.

Python utility

Finally, a python utility *db_populate.py* that would:

- scrape .csv files in directory /raw_data
- do POST API method calls using *Requests* library to INSERT data into corresponding tables
- (AND) do GET API method calls to verify object counts

The result of the GET API method is a JSON body comprising all entries in the table and a *result_count* indicating number of table entries.

Here's an excerpt of the python utility:

```
5
6 with open('raw_data/%s' %FILE) as f:
7     lines = [line.rstrip() for line in f]
8
9 path = {'prod': '/api/v1/production/',
10        'nfs_info': '/api/v1/nfs_info/',
11        'server': '/api/v1/server/',
12        'server_make': '/api/v1/server_make/',
13        'fault_ticket': '/api/v1/fault_ticket/',
14        'switch': '/api/v1/switch/',
15        'vlan_mapping': '/api/v1/vlan_mapping/',
16        'vlan': '/api/v1/vlan/',
17        'switch_port_mapping': '/api/v1/switch_port_mapping/',
18        'server_switch_info': '/api/v1/server_switch_info/'
19    }
20
21 header = True
22 url = endpoint + path[FILE]
23 for line in lines:
24     if header:
25         headers = line.split()
26         header = False
27         continue
28     data = {}
29     index = 0
30     entry = line.split()
31     while index < len(entry):
32         data[headers[index]] = entry[index]
33         index += 1
34     r = requests.post(url = url, json = data)
35     print(r)
36 g = requests.get(url = url)
37 print(g.text)
```

Figure-7: *db_populate.py* utility to populate tables

Forking function – Bridge between API and database

Function *common_db_crud_ops* forks incoming API requests to corresponding database operations (to be extended to PATCH and DELETE)

```
434
435 def common_db_crud_ops(table, request):
436     if request.method == 'POST':
437         if request.is_json:
438             status = db_testinfra.insert(table, request.json)
439             return status
440     if request.method == 'GET':
441         data = db_testinfra.get(table)
442         return jsonify(data)
```

Figure-8: Function forking api calls to corresponding DB ops

SQL – Under the hood data persistence

And finally, it is down to the select few functions to insert and fetch data to/from the database and thus persist it. Below are the *insert* and *get* methods implemented within the *SqlLibTestInfra* class:

```
107
108 def insert(self, table, data_in_json):
109     if table == 'nfs_info':
110         _nfs_id = data_in_json['_nfs_id']
111         _nfs_ip_addr = data_in_json['_nfs_ip_addr']
112         try:
113             self.conn.execute('insert into nfs_info (nfs_id, nfs_ip_addr) \
114                               VALUES(%d, "%s");' % (_nfs_id, _nfs_ip_addr))
115             self.conn.commit()
116             return (return_codes['insert_success'])[1],
117                    return_codes['insert_success'][0]
118         except Error as e:
119             return (str(e) + '\n', 500)
120
121     if table == 'production':
122         _prod_id = int(data_in_json['production_id'])
123         _prod_name = data_in_json['production_name']
124         _nfs_share_fqdn = data_in_json['nfs_share_fqdn']
125         _outer_vc_ip = data_in_json['outer_vc_ip']
126         _inner_vc_ip = data_in_json['inner_vc_ip']
127         try:
128             _desc = data_in_json['description']
129             _desc = ''
130         except KeyError:
131             _runner_ip = data_in_json['runner_ip']
132             _nfs_id = int(data_in_json['_nfs_id'])
133             try:
134                 self.conn.execute('insert into production (production_id, \
135                               production_name, nfs_share_fqdn, outer_vc_ip, \
136                               inner_vc_ip, description, runner_ip, nfs_id) \
137                               VALUES(%d, "%s", "%s", "%s", "%s", "%s", "%s", "%s");' \
138                               % (_prod_id, _prod_name, _nfs_share_fqdn, _outer_vc_ip,
139                               _inner_vc_ip, _desc, _runner_ip, _nfs_id))
140                 self.conn.commit()
141                 return (return_codes['insert_success'])[1],
142                        return_codes['insert_success'][0]
143             except Error as e:
144                 return (str(e) + '\n', 500)
145
288
289 def get(self, table):
290     self.conn.row_factory = sqlite3.Row
291     if table == 'nfs_info':
292         rows = self.conn.execute('select * from nfs_info;')
293     elif table == 'production':
294         rows = self.conn.execute('select * from production;')
295     elif table == 'server':
296         rows = self.conn.execute('select * from server;')
297     elif table == 'server_make':
298         rows = self.conn.execute('select * from server_make;')
299     elif table == 'fault_ticket':
300         rows = self.conn.execute('select * from fault_ticket;')
301     elif table == 'switch_port_mapping':
302         rows = self.conn.execute('select * from switch_port_mapping;')
303     elif table == 'switch':
304         rows = self.conn.execute('select * from switch;')
305     elif table == 'server_switch_info':
306         rows = self.conn.execute('select * from server_switch_info;')
307     elif table == 'vlan':
308         rows = self.conn.execute('select * from vlan;')
309     elif table == 'vlan_mapping':
310         rows = self.conn.execute('select * from vlan_mapping;')
311     _res = OrderedDict()
312     _res_list = []
313     for row in rows.fetchall():
314         _res_list.append(dict(row))
315     _res['result_count'] = len(_res_list)
316     _res['results'] = _res_list
317     return _res
318
```

Figure-9: *insert* function to INSERT data into database AND *get* function to SELECT data and return in key-value/JSON format

Analytics

The idea with this framework was to make it efficient for Developers to consume resources within their prod. environments by inheriting APIs outlined in this document (including but not limited to, yet).

Subsequently, it is also critical to capture information and monitor resources given the scale spans several servers and switches. After all, 'No resource pool is infinite' AND '..everybody wants resources!'. While the dashboard will significantly evolve in due course of time this section attempts to capture in gist, the naïve first few steps.

SQL – Beyond native CRUD

Putting together some explorations with SQL statements (comprising JOINS, GROUP BY etc.) that helped gain insights to set the stage for rich analytics.

Here are a few initial metrics to visualize:

- Server usage – grouped by each prod./Dev environment
- Switch usage – group by each prod./Dev environment
- Server usage – grouped by make, memory & cpu specs.

Below are corresponding SQL queries:

```
318
319 def custom_queries(self, query_type):
320     if query_type == 'GET_SERVERS_PER_PROD':
321         _res = self.conn.execute('select production.production_name, ' \
322                                 'count(server.server_name) from server ' \
323                                 'JOIN production ON ' \
324                                 'server.production_id = production.production_id ' \
325                                 'GROUP BY production.production_name')
326     if query_type == 'GET_SWITCHES_PER_PROD':
327         _res = self.conn.execute('select production.production_name, ' \
328                                 'count(server_switch_info.switch_port_id) ' \
329                                 'from server ' \
330                                 'JOIN production ON ' \
331                                 'server.production_id = production.production_id ' \
332                                 'JOIN server_switch_info ON ' \
333                                 'server.server_name = server_switch_info.server_name ' \
334                                 'JOIN switch_port_mapping ON ' \
335                                 'server_switch_info.switch_port_id = switch_port_mapping.switch_port_id ' \
336                                 'GROUP BY production.production_name')
337     if query_type == 'GET_SERVER_TYPES':
338         _res = self.conn.execute('select server_make.manufacturer || server_make.memory_in_gb || server_make.cpu_cores AS server_type, ' \
339                                 'count(server_make.manufacturer) from server ' \
340                                 'JOIN server_make on server.server_make_id = server_make.server_make_id ' \
341                                 'GROUP BY (server_type)')
342     return _res.fetchall()
```

Figure-10: Queries to generate data insights

Dashboard

Data from *custom_queries* is packaged as labels and values using the python zip function before using Flask's *render_template* utility:

```
+ 537     return render_template('chart.html',
+ 538                             title_1='SERVERS - Grouped by Production',
+ 539                             max_1=int(1.1 * max(values_1)),
+ 540                             labels_1=labels_1,
+ 541                             values_1=values_1,
+ 542                             title_2='SWITCHES - Grouped by Production',
+ 543                             max_2=int(1.1 * max(values_2)),
+ 544                             labels_2=labels_2,
+ 545                             values_2=values_2,
+ 546                             title_3='SERVERS - Grouped by H/W Specs',
+ 547                             max_3=int(1.1 * max(values_3)),
+ 548                             set=zip(values_3, labels_3, colors)
+ 549                             )
```

Figure-11: Queries to generate data insights

The html code in *chart.html* translates the data into bar and pie charts.

Below is a first peek at the dashboard accessed via a browser:

URL <http://<ip-address>:5000/>




Figure-12: Infrastructure analytics dashboard

Conclusion and Future work

This project has long been on my mind and I am grateful it is finally underway (and beyond) during the course of this semester.

Achievements include:

- Conceptualization to realization using techniques imbibed via IST-659 in designing databases
 - All tables organically went from ERD through logical-modeling and normalization process before being 3NF compliant
- All tables have real data from THREE of our prod. environments which greatly establishes a workflow we have in place for consumption
- Setting in place a rich web application framework

Future work

There is significant work lined up and I have captured a few:

1. Implementation of PATCH and DELETE API methods that correspondingly translate to SQL UPDATE and DELETE operations. This is an essential part of the CRUD operations that is pending implementation.

Since server allocations can constantly change hands across prod. environments it is imperative these methods be available before rolling the framework out to Developer community.

2. Use concepts of IST-659:
 - Use of procedures/functions to reduce error and code-redundancy
 - Efficient SQL commands
3. Migrating data for several hundred servers and switches across prod. environments to the web framework and database.

Subsequently, programming switches with network configuration using Tcl/Tk, expect or ansible seems a tangible option since all the data is co-located and available across a few tables in the database.

4. Extend analytics dashboard to comprise more data and visualizations – work with our internal stakeholders regarding critical metrics.

Appendix

Front-end code for dashboard (chart.html)

Below is the code under /templates, namely chart.html which is used by Flask to route GUI requests via the render_template utility

```

1 <!DOCTYPE html>
2 <html lang="en">
3 <head>
4 <meta charset="utf-8" />
5 <title>{{ title }}</title>
6 <script src="https://cdnjs.cloudflare.com/ajax/libs/Chart.js/1.0.2/Chart.min.js"></script>
7 </head>
8
9 <body>
10 <style>
11 h2 {
12   position: absolute;
13   left: 100px;
14   top: 0px;
15 }
16 </style>
17 <h2 style="font-family:impact; font-size:15px">{{ title_1 }}</h2>
18 <canvas id="chart-1" width="400" height="300" \
19 left="100px", style="left:100px; top: 70px; position: absolute;"></canvas>
20 <script>
21 var options = {
22   chartArea: {left:0, top:0}
23 }
24 // bar chart data
25 var barData = {
26   labels : [
27     {% for item in labels_1 %}
28     '{{ item }}',
29     {% endfor %}
30 ],
31   datasets : [
32     {
33       fillColor: "rgb(151,187,206,0.2)",
34       strokeColor: "rgb(151,187,206,1)",
35       pointColor: "rgb(151,187,206,1)",
36       data : [
37         {% for item in values_1 %}
38         '{{ item }}',
39         {% endfor %}
40 ]
41 }
42 ]
43 }
44
45 // get bar chart canvas
46 var mychart = document.getElementById("chart-1").getContext("2d");
47
48 steps = 10
49 max = (max_1)
50
51 // draw bar chart
52 new Chart(mychart).Bar(barData, {
53   scaleOverride: true,
54   scaleSteps: steps,
55   scaleStepwidth: Math.ceil(max / steps),
56   scaleStartValue: 0,
57   scaleShowVerticalLines: true,
58   scaleShowGridLines: true,
59   barShowStroke: true,
60   scaleShowLabels: true
61 }
62 );
63
64
65
66
67
68
69
70
71
72
73 <body>
74 <head>
75 <style>
76 h4 {
77   position: absolute;
78   left: 1100px;
79   top: 0px;
80 }
81 </style>
82 </head>
83 <h4 style="font-family:impact; font-size:15px">{{ title_3 }}</h4>
84 <canvas id="chart-3" width="400" height="300" \
85 left="100px", style="left:1000px; top: 70px; position: absolute;"></canvas>
86 <script>
87 var pieData = [
88   {% for item, label, colors in set %}
89   {
90     value: {{item}},
91     label: "{{label}}",
92     color : "{{colors}}"
93   },
94   {% endfor %}
95 ];
96
97 // get bar chart canvas
98 var mychart = document.getElementById("chart-3").getContext("2d");
99 steps = 10
100 max = (max_3)
101
102 // draw pie chart
103 new Chart(document.getElementById("chart-3").getContext("2d")).Pie(pieData)
104
105 </script>
106 </body>
107 </html>

```

Figure-13: HTML code to create a bar-chart and pie-chart